

Statistical Machine Learning: Assignment 2

Joris van Vugt, s4279859

October 19, 2016

Exercise 1 – Sequential learning

Part 1 – Obtaining the prior

1. First we compute the precision matrix using `numpy.linalg.inv`:

$$\tilde{\Lambda} = \tilde{\Sigma}^{-1} = \left(\begin{array}{cc|cc} 60 & 50 & -48 & 38 \\ 50 & 50 & -50 & 40 \\ \hline -48 & -50 & 52.4 & -41.4 \\ 38 & 40 & -41.4 & 33.4 \end{array} \right) = \left(\begin{array}{c|c} \tilde{\Lambda}_{aa} & \tilde{\Lambda}_{ab} \\ \hline \tilde{\Lambda}_{ba} & \tilde{\Lambda}_{bb} \end{array} \right)$$

Now, using equations 2.73 and 2.75 from Bishop, we can compute the conditional covariance:

$$\Sigma_p = \tilde{\Lambda}_{aa}^{-1} = \begin{pmatrix} 0.1 & -0.1 \\ -0.1 & 0.12 \end{pmatrix}$$

and the conditional mean:

$$\begin{aligned} \mu_p &= \mu_a - \tilde{\Lambda}_{aa}^{-1} \tilde{\Lambda}_{ab} (x_b - \mu_b) \\ &= \begin{pmatrix} 1 \\ 0 \end{pmatrix} - \begin{pmatrix} 0.1 & -0.1 \\ -0.1 & 0.12 \end{pmatrix} \begin{pmatrix} -48 & 38 \\ -50 & 40 \end{pmatrix} \left(\begin{pmatrix} 0 \\ 0 \end{pmatrix} - \begin{pmatrix} 1 \\ 2 \end{pmatrix} \right) \\ &= \begin{pmatrix} 1 \\ 0 \end{pmatrix} - \begin{pmatrix} 0.1 & -0.1 \\ -0.1 & 0.12 \end{pmatrix} \begin{pmatrix} -48 & 38 \\ -50 & 40 \end{pmatrix} \begin{pmatrix} -1 \\ -2 \end{pmatrix} \\ &= \begin{pmatrix} 1 \\ 0 \end{pmatrix} - \begin{pmatrix} 0.2 & -0.2 \\ -1.2 & 1 \end{pmatrix} \begin{pmatrix} -1 \\ -2 \end{pmatrix} \\ &= \begin{pmatrix} 1 \\ 0 \end{pmatrix} - \begin{pmatrix} 0.2 \\ -0.8 \end{pmatrix} \\ &= \begin{pmatrix} 0.8 \\ 0.8 \end{pmatrix} \end{aligned}$$

- 2.

```
def generate_pair():  
    return np.random.multivariate_normal([0.8, 0.8],  
                                          [[0.1, -0.1],  
                                           [-0.1, 0.12]])
```

$$\mu_t = \begin{pmatrix} 0.28 \\ 1.18 \end{pmatrix}$$

3. To calculate the probability density of our multivariate Gaussian random variable, we use `scipy.stats.multivariate_normal` and its `pdf` method.




Figure 1: Probability density plot of the multivariate Gaussian

Part 2 – Generating the data

```
1. N = 1000
   data = np.random.multivariate_normal([0.28, 1.18],
                                         [[2.0, 0.8],
                                          [0.8, 4.0]],
                                         N)

   np.savetxt('data.txt', data)
```

2. The maximum likelihood estimate of the mean is simply the mean of the observed data:

$$\boldsymbol{\mu}_{\text{ML}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n = \begin{pmatrix} 0.25 \\ 1.21 \end{pmatrix}$$

Computing the maximum likelihood estimate of the covariance is slightly more involved:

$$\boldsymbol{\Sigma}_{\text{ML}} = \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \boldsymbol{\mu}_{\text{ML}})(\mathbf{x}_n - \boldsymbol{\mu}_{\text{ML}})^T = \begin{pmatrix} 2.023 & 0.828 \\ 0.828 & 3.626 \end{pmatrix}$$

To compute the unbiased maximum likelihood covariance estimate, we normalize by $N - 1$ instead of N :

$$\boldsymbol{\Sigma}_{\text{ML}} = \frac{1}{N-1} \sum_{n=1}^N (\mathbf{x}_n - \boldsymbol{\mu}_{\text{ML}})(\mathbf{x}_n - \boldsymbol{\mu}_{\text{ML}})^T = \begin{pmatrix} 2.025 & 0.829 \\ 0.829 & 3.629 \end{pmatrix}$$

These results were obtained with the following code:

```
mu_ml = data.mean(axis=0)
x = data - mu_ml
cov_ml = np.dot(x.T, x) / N
cov_ml_unbiased = np.dot(x.T, x) / (N - 1)
```

Note that the left factor is transposed instead of the right factor in the covariance estimates. This is because our points are row vectors instead of column vectors (i.e., `data` has shape $N \times 2$).

We can compare our estimates to the true statistics:

$$\boldsymbol{\mu}_t = \begin{pmatrix} 0.28 \\ 1.18 \end{pmatrix} \quad \boldsymbol{\Sigma}_t = \begin{pmatrix} 2.0 & 0.8 \\ 0.8 & 4.0 \end{pmatrix}$$

Our estimates are pretty close to their true values. The unbiased estimate of the covariance is not much closer than the biased estimate. Because N is relatively high, the slight change in normalization does not have much effect.

Part 3 – Sequential learning algorithms

1. Using equation 2.126 from Bishop

$$\boldsymbol{\mu}_{\text{ML}}^{(N)} = \boldsymbol{\mu}_{\text{ML}}^{(N-1)} + \frac{1}{N}(\mathbf{x}_N - \boldsymbol{\mu}_{\text{ML}}^{(N-1)})$$

we can come up with a Python procedure for sequential learning:

```
mu_ml = 0
for i in range(N):
    mu_ml += (data[i]-mu_ml) / (i+1)
```

The starting value $\boldsymbol{\mu}_{\text{ML}}^{(0)}$ does not matter, but is arbitrarily set to 0. We divide by $i + 1$ instead of just i , because Python uses 0-indexing.

- 2.